Doomscroll -- Lab 3 Dialog Map Refinement

Per Lab 3 Section 3.1.4 ("re-visit your dialog map from Lab #2, add detail"). Diagram carried forward unchanged: dialog_map_lab2.jpg / .pdf (7 states, 15 transitions -- see dialog_map.pdf for the full transition legend table).

Why no structural changes were made

The 7 states and 15 transitions are still the correct target UI flow -- nothing about the app's intended behavior changed since Lab 2. What changed is implementation status, not the dialog map itself: none of the 7 states exist as real UI screens yet (see design_model.md Section 1 -- no GUI framework has been adopted). Retrofitting the diagram to show "what cv2 window is currently on screen" would conflate a design artifact with an implementation-status note, so that status is tracked separately here instead of by editing the diagram.

Implementation-status mapping (new for Lab 3)

Dialog map state | Status
Setup (`s1`) | `[PLANNED]` -- no UI; today, camera access is implicit (first `cv2.VideoCapture` call)
Calibration (`s2`) | `[BUILT, different UI]` -- `scripts/calibrate.py` / `scripts/quick_calibrate.py` run a real calibration point sequence, but as a full-screen `cv2` window, not the mockup's dialog-based UI
Reading -- Autoscroll Disabled | `[PLANNED]` -- no concept of "disabled vs. active" state exists; today the app is either running the live loop or not running at all
Reading -- Autoscroll Active (`s4`) | `[BUILT, different UI]` -- the live loop in `scripts/phase2_gaze_scroll.py`, with the `cv2` debug overlay standing in for `ReaderOverlay`
Reading -- Autoscroll Paused | `[BUILT]` -- `AttentionGate` already implements a form of this (pauses the dwell trigger when the user looks away), though it is not user-initiated as transition 8 specifies
Recalibration Prompt (`s5`) | `[PLANNED]` -- no confidence-based recalibration prompt exists; `AttentionGate`'s pause is the closest built analogue
Settings (`s3`) | `[PLANNED]` -- no settings UI; current equivalent is re-running the script with different CLI flags (`--label-mode`, `--lag-ms`, etc. for training; no live-app settings at all)

One real behavioral difference worth flagging

Transition 5 (Reading -- Autoscroll Active -> Recalibration Prompt, "sustained low confidence") assumes a single, static notion of "confidence." In practice (see troubleshooting during Lab 3 testing -- the ML model's CLAMP_FRAC extrapolation behavior under out-of-distribution head poses), confidence degradation is more nuanced than a single scalar threshold: a trained-model GazePredictionStrategy can produce confidently-wrong predictions rather than failing loudly. This is noted as a design risk to revisit when RecalibrationPromptDialog is actually built -- the trigger condition for transition 5 may need to be strengthened beyond "confidence below threshold" (e.g., also tracking prediction variance/instability over a short window).